

Temporal segmentation of long videos is an important problem, that has largely been tackled through supervised learning, often requiring large amounts of anno