

Highlight detection has the potential to significantly ease video browsing, but existing methods often suffer from expensive supervision requirements, where human